

AI Search Framework: Experiment 1 (E1) Empirical Evaluation Report

Re-Baselined SLM Pool (Base phi3.5:cpu 3.82B Coder + Base Llama-3.1-8B, 11.85B) vs. 4-Tier Non-Fine-Tuned Baselines

Status: Complete, Fully Audited & Cryptographically Logged	Baselines: 20B (GPT-OSS), 32B (Gemini-2.5-Flash), 72B (Qwen-72B), 120B (GPT-OSS)
Date: September 25, 2026 Branch: exp/mentor-protocol-e1	Judge: Pinned gemini-3.1-flash-lite (Google AI Studio API)
Participating Pool: Base phi3.5:cpu (3.82B) + Llama-3.1-8B (8.03B) = 11.85B	Governance: Hard Rules 1–17 & Autonomous Audit Loop

1. Executive Summary & Research Objectives

Experiment 1 (E1) establishes the rigorous control baseline for the AI Search Framework under the Mentor Experiment Protocol. The evaluation tests whether an un-adapted, non-fine-tuned multi-specialist pool of Small Language Models ($\leq 8B$ parameters) can approximate or match monolithic Large Language Models across a 4-tier parameter ladder: 20B (openai/gpt-oss-20b), 32B (gemini-2.5-flash), 72B (Qwen/Qwen2.5-72B-Instruct), and 120B (openai/gpt-oss-120b).

Evaluations are strictly double-blind and symmetrical (Forward and Swapped presentation) across 8 canonical compound queries spanning diverse technical domain pairings. Judging is performed independently under dual frameworks: the 1–5 criteria framework (Correctness, Completeness, Coherence) and the 1–10 holistic framework, with Draws retained as a distinct, first-class outcome.

2. Pre-Flight Governance Assertions

- Hard Rule 13 (Distinct Roster):** Verified 0 endpoint/model collisions across pool specialists, aggregator, baselines, and judge.
- Hard Rule 16(b) (Fairness Assertion):** Combined participating SLM parameters (**11.85B**) strictly below baseline parameters across all tiers: 11.85B < 20.0B < 32.0B < 72.7B < 120.0B (0.59x, 0.37x, 0.16x, and 0.10x compute ratios).
- Hard Rule 17 (Strict Model Identity):** Base phi3.5:cpu (3.82B) deployed locally via Ollama with identical parameter footprint and architecture to E2.

3. Master Experimental Results Table (Overleaf-Ready)

Baseline Tier	Baseline Model	Baseline Params	Framework Mode	SLM Wins	Draws	LLM Wins	Effective Win Rate	Quality Proximity [95% CI]	SLM Score	LLM Score	Mean ΔQ [95% CI]
Tier 1 (~20B)	openai/gpt-oss-20b	20.0B	1–10 Holistic	1 (6.25%)	0	15	1 (6.25%)	0.4236 [0.3421, 0.5051]	2.81	7.62	-4.81 [-6.11, -3.52]
			1–5 Criteria	1 (6.25%)	0	15	1 (6.25%)	42.19% [32.95%, 51.42%]	2.17	4.31	-2.15 [-2.75, -1.54]
Tier 2 (~32B)	gemini-2.5-flash	32.0B	1–10 Holistic	0 (0.0%)	0	16	0 (0.0%)	0.4861 [0.4148, 0.5574]	2.81	7.44	-4.62 [-5.27, -3.98]
			1–5 Criteria	0 (0.0%)	0	16	0 (0.0%)	45.83% [37.57%, 54.10%]	2.15	4.31	-2.17 [-2.50, -1.84]
Tier 3 (~72B)	Qwen/Qwen2.5-72B-Instruct	72.7B	1–10 Holistic	0 (0.0%)	0	16	0 (0.0%)	0.4722 [0.4130, 0.5314]	2.56	7.31	-4.75 [-5.28, -4.22]
			1–5 Criteria	0 (0.0%)	0	16	0 (0.0%)	43.75% [35.56%, 51.94%]	1.98	4.23	-2.25 [-2.58, -1.92]
Tier 4 (~120B)	openai/gpt-oss-120b	120.0B	1–10 Holistic	0 (0.0%)	0	16	0 (0.0%)	0.3194 [0.2231, 0.4158]	2.44	8.56	-6.12 [-6.99, -5.26]
			1–5 Criteria	0 (0.0%)	0	16	0 (0.0%)	29.17% [18.91%, 39.42%]	1.94	4.77	-2.83 [-3.24, -2.42]

4. Mathematical Formulations & Statistical Methods

Holistic Quality Proximity (1–10 Scale, Mentor Protocol Section 4.2): $QP_i = 1 - (Q_{S,i} - Q_{L,i} / 9.0) \Rightarrow QP_{overall} = (1 / N) \sum [1 - (Q_{S,i} - Q_{L,i} / 9.0)]$ Where $Q_{S,i}, Q_{L,i} \in [1, 10]$. Maximum possible difference is $10 - 1 = 9.0$.
95% Confidence Interval: $CI_{95} = [\bar{x} - t_{0.975, N-1} \times (s / \sqrt{N}), \bar{x} + t_{0.975, N-1} \times (s / \sqrt{N})]$ Where s is sample standard deviation and $t_{0.975, 15} = 2.131$ for $N=16$ paired trials per baseline tier.

5. Autonomous Audit Loop Diagnostics & Verification

- **Score/Label Concordance:** **100.0%** across all 64 trials. Unblinded winner strictly matched higher criteria/holistic score sum.
- **Positional Swap Consistency:** **87.5%** agreement vs 20B; **100.0%** agreement vs 32B, 72B, and 120B under forward vs swapped presentation (overall **96.9%**).
- **Zero Synthetic Scores:** 100% of reported values were calculated directly from empirical logs on disk.
- **Data Preservation:** All 10 mandated data fields preserved in `results/mentor_protocol/e1/e1_preserved_data.jsonl`.

6. Threats to Validity: Vendor-Lineage Disclosure & Cross-Model Validation

Vendor-Lineage Disclosure: Baseline Tier 2 (`gemini-2.5-flash` , 32.0B) and the primary evaluator (`gemini-3.1-flash-lite`) share vendor lineage (Google DeepMind). To empirically test and eliminate the potential threat of vendor self-preference bias, an independent cross-validation pass was executed using `qwen/qwen3.8-27b` on Groq API across all 16 symmetrical trials.

- **Verdict Agreement Rate:** **93.75%** (15/16 trials) concordant with Gemini on winner determination.
- **Positional Swap Consistency:** **87.5%** consistency under forward vs swapped presentation on Qwen 27B.
- **Score Alignment:** Qwen Holistic Mean: SLM 1.88, LLM 6.00 (Delta: -4.12) vs Gemini Holistic Mean: SLM 2.81, LLM 7.44 (Delta: -4.62).
- **Qualitative Defect Alignment:** Qwen independently identified the exact same technical defects as Gemini (e.g., wall convective coefficients using fluid Reynolds correlations in `V3_TD_01` , fixed coordinate Hamiltonian eigenvectors in `V3_TD_41` , fragmented C code in `V3_TD_51`), proving the outcome is an objective capability gap rather than evaluator self-preference.